

Object Detection and Faster RCNN

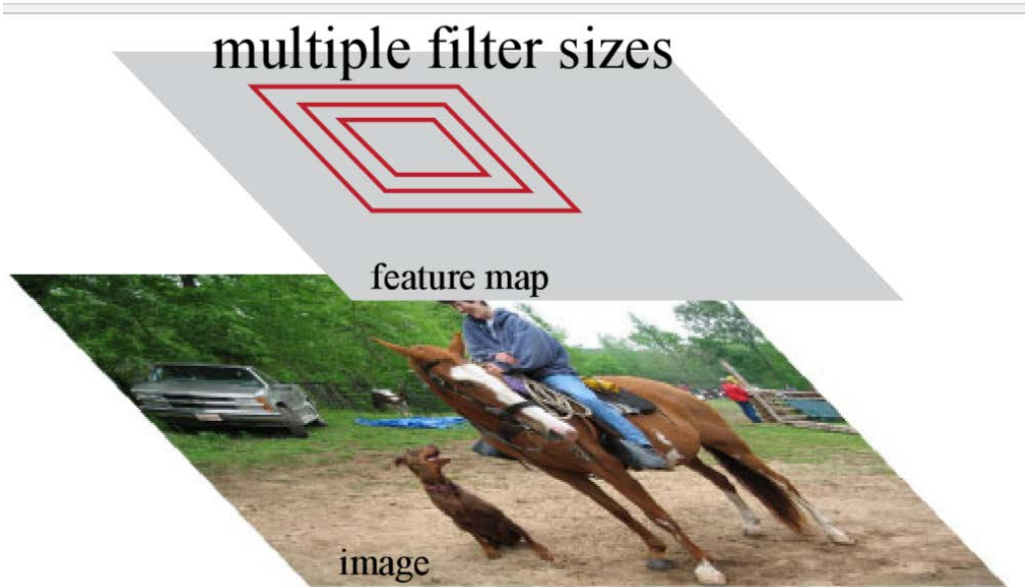
Paul Rodriguez SDSC



How to find object instances of different sizes across images?



Combination of image resizing, different filter sizes, and sliding windows helps find objects




Combination of image resizing, different filter sizes, and sliding windows helps find objects



Or, use CNN to classify and detect: Faster RCNN


(Ren,He,Girschick,Sun)

**A side
path to
select
Regions**



Or, use CNN to classify and detect: Faster RCNN


(Ren,He,Girschick,Sun)



*Images are normalized
and resized*

Or, use CNN to classify and detect: Faster RCNN

(Ren,He,Girschick,Sun)




Pool with adjustable windows to transform into one size vector for classification

Images are normalized and resized

Or, use CNN to classify and detect: Faster RCNN

(Ren,He,Girschick,Sun)

1. Start with 2000 sampled regions; segment, group, get texture of possible foreground regions
2. Then use feature map values and a model to predict if an object (of any class) would be detected in each window.
3. Pass maps for best regions to classifier




Pool with adjustable windows to transform into one size vector for classification

Images are normalized and resized

Training Data

- Given positive sample, generate negative samples (and balance sample sizes)



Sample and add boxes with $< 40\%$ overlap as 'negative', $> 50\%$ as positive

Training Data

- Given positive sample, generate negative samples (and balance numbers)

Train region proposal and classifier alternatively in 4 phases

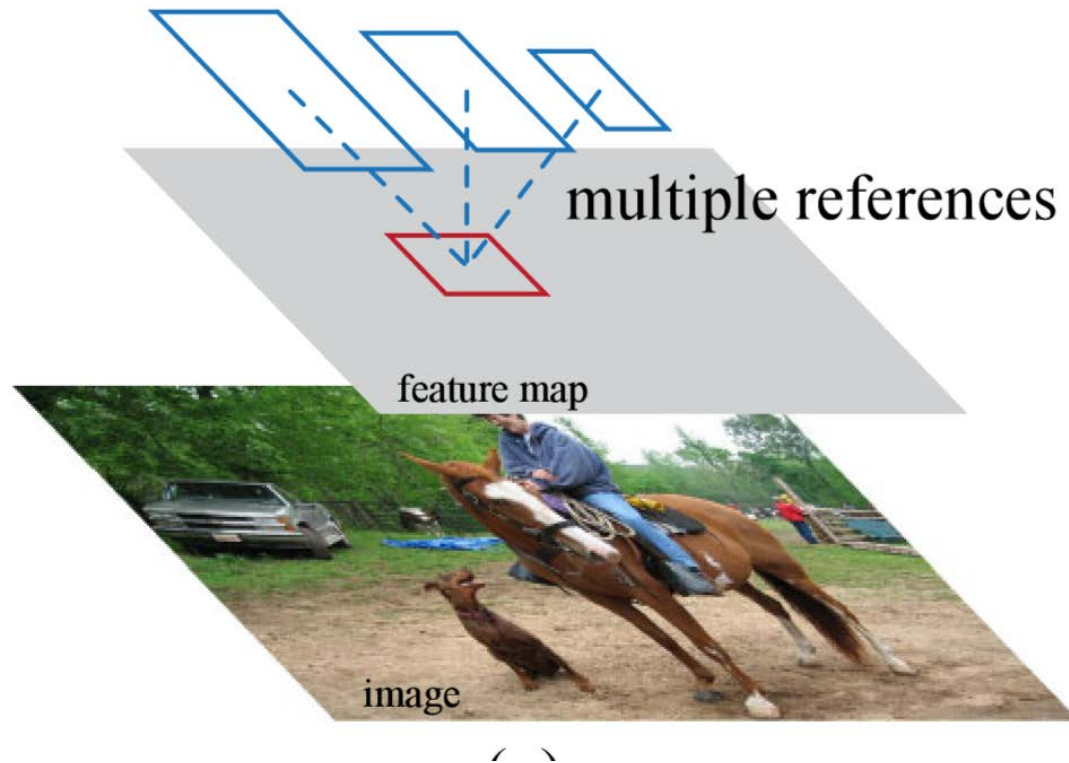
Phases:

1,3 region proposal

2,4 classification

Region output

- Output bounding box information (box center, height, width)



Region Box size ranges

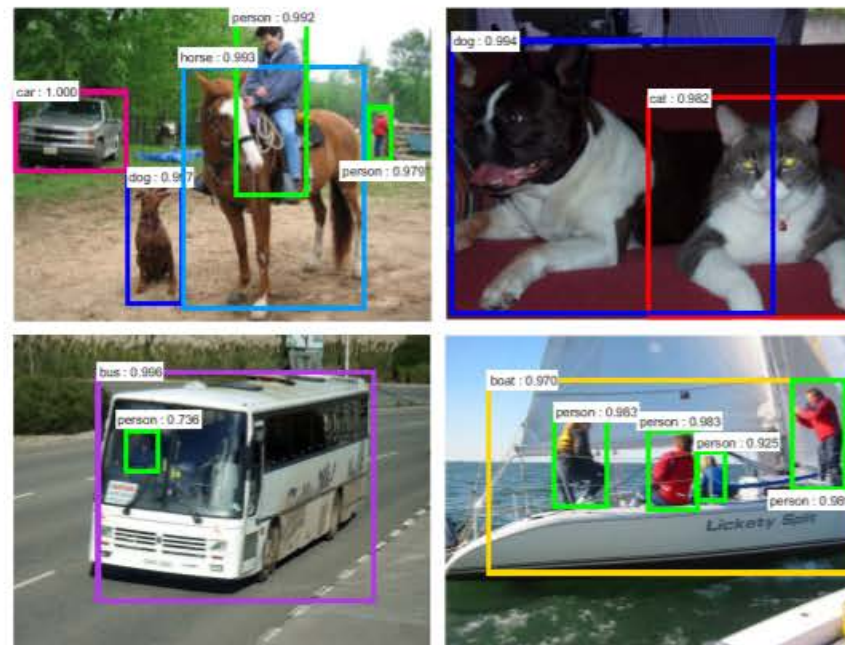
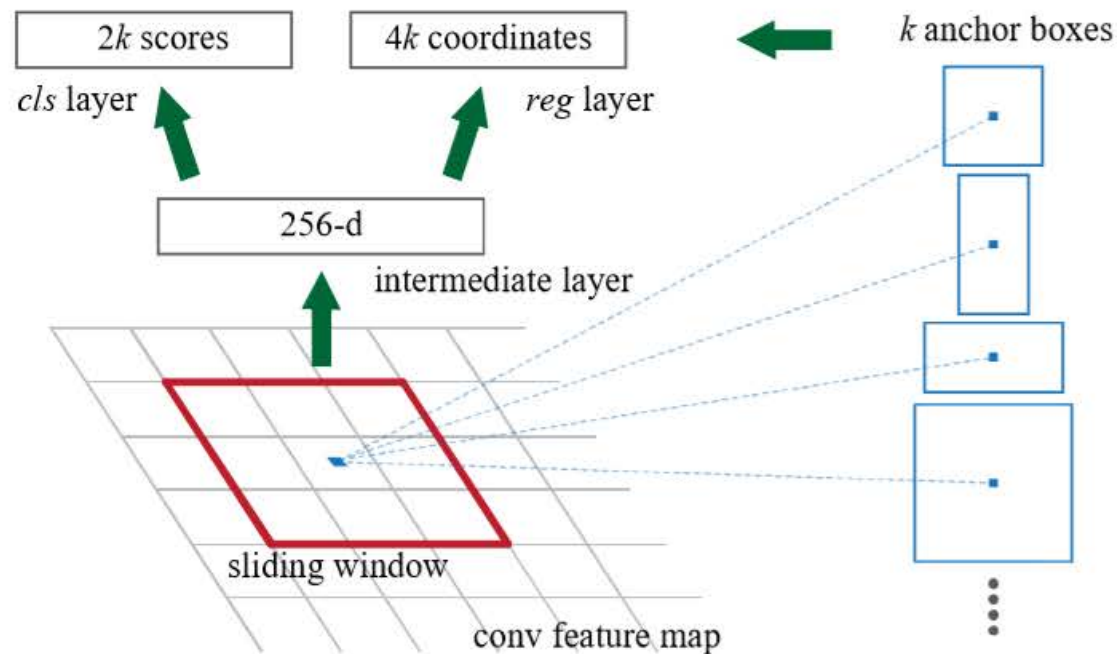




Figure 3: **Left:** Region Proposal Network (RPN). **Right:** Example detections using RPN proposals on PASCAL VOC 2007 test. Our method detects objects in a wide range of scales and aspect ratios.

Samples of 'hand' training data

Person boxes (using YOLO in
'darknet' code)

Face detection using 'DLIB' library

Hand boxes were hand drawn(Mittal etal.)



Matlab Implementation

```
detector = trainFasterRCNNObjectDetector(  
    trainingData, layers, options, ...
```


Table of ~4K image file
names and ~8K boxes



A defined set of layers



learning rate, epochs
.0001 phase 1,2
.00001 phase 3,4
15 epochs



Matlab Implementation


```
detector = trainFasterRCNNObjectDetector(  
    trainingData, layers, options, ...  
    'NegativeOverlapRange', [0.1 0.4], ...  
    'PositiveOverlapRange', [0.6 1], ...  
    'NumStrongestRegions', 200, ...  
    'BoxPyramidScale', 1.2, ...  
    'NumBoxPyramidLevels', 5);
```

For making pos/neg samples

For selecting object (anchor) boxes

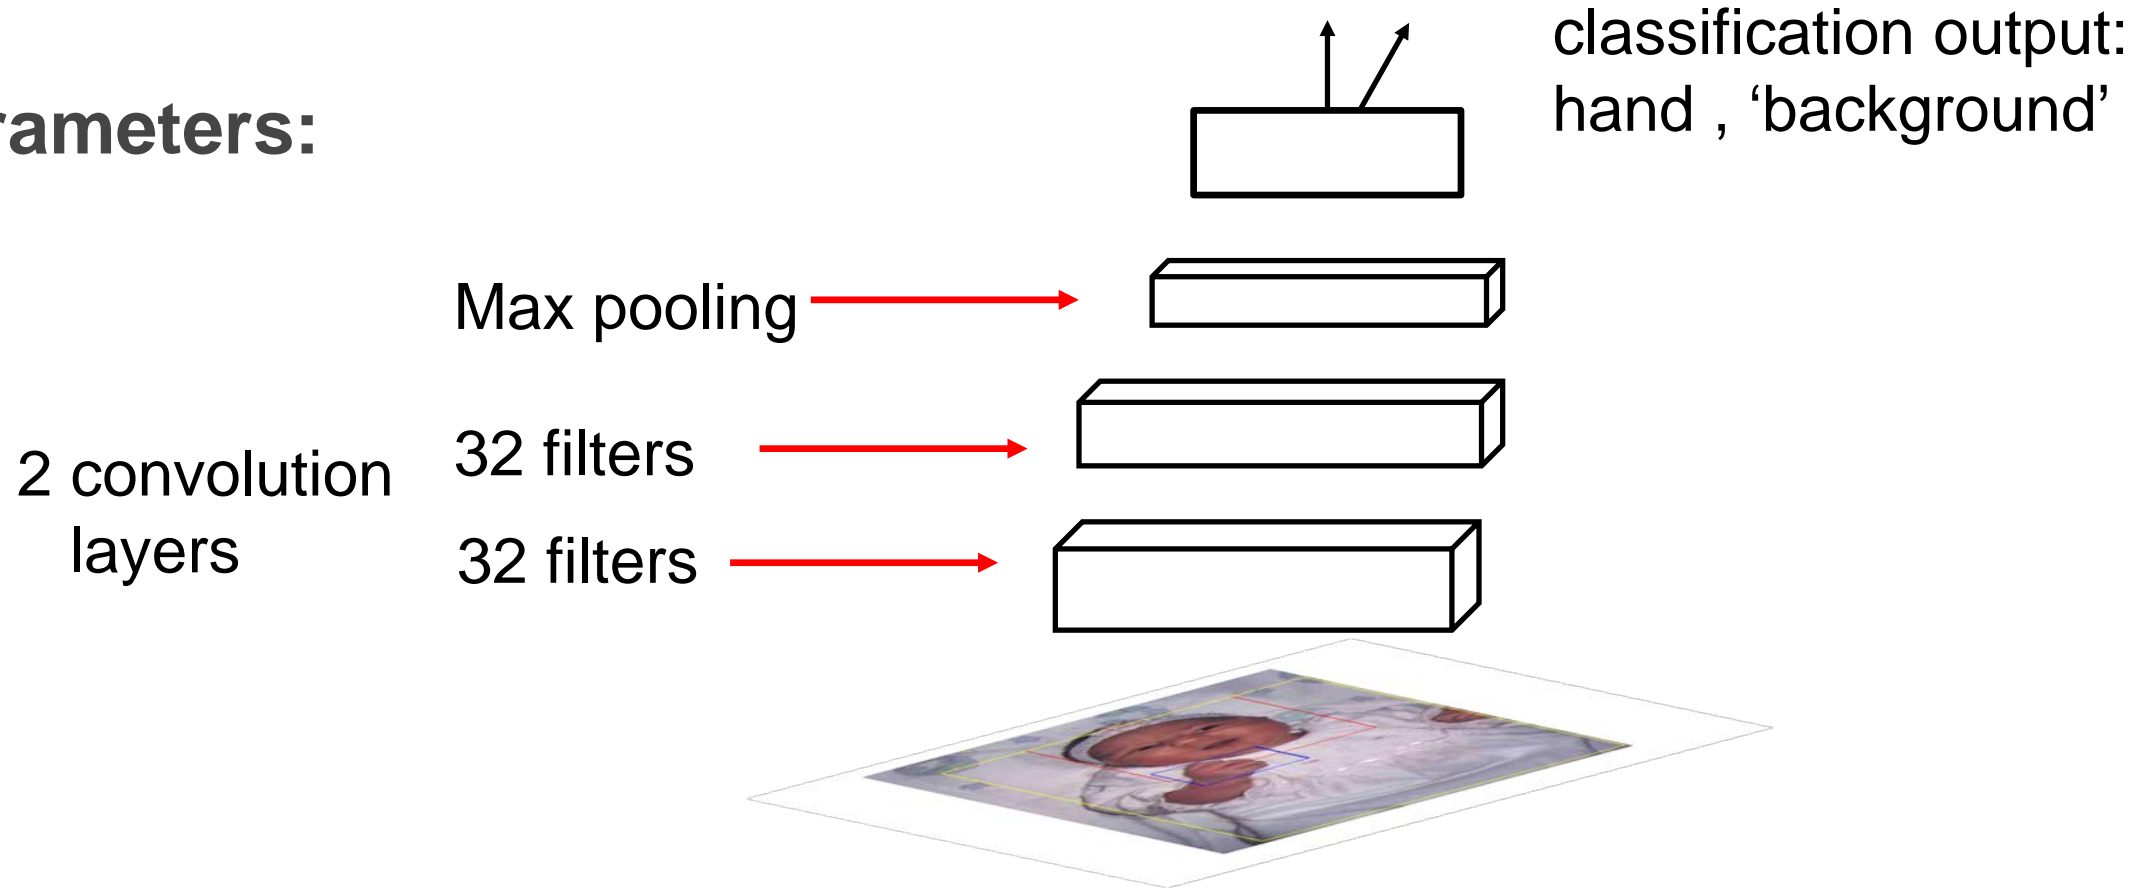
Matlab Faster RCNN

- My parameters:



Matlab Faster RCNN

- **My parameters:**



Network detail



>> detector.Network.Layers

11x1 Layer array with layers:

1	'imageinput'	Image Input	128x128x3 images with 'zerocenter' normalization
2	'conv_1'	Convolution	32 5x5x3 convolutions with stride [2 2] and padding [1 1 1 1]
3	'relu_1'	ReLU	ReLU
4	'conv_2'	Convolution	32 5x5x32 convolutions with stride [2 2] and padding [1 1 1 1]
5	'relu_2'	ReLU	ReLU
6	'roi pooling layer'	ROI Max Pooling	ROI Max Pooling with grid size [14 14]
7	'fc_1'	Fully Connected	32 fully connected layer
8	'relu_3'	ReLU	ReLU
9	'fc_2'	Fully Connected	2 fully connected layer
10	'softmax'	Softmax	softmax
11	'classoutput'	Classification Output	crossentropyex with classes 'hand' and 'Background'

Results (so far)

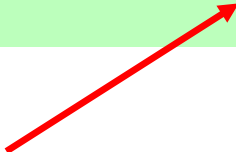
- Using 50% overlap with true box as correct ~25% TP rate
- 4-8 hours on 1 compute node (CPU) 15 epochs ~4K images



Matlab Implementation

```
detector = trainFasterRCNNObjectDetector(  
    trainingData, vgg19, options, ...
```


Table of ~1K image file
names and ~2K boxes



A pretrained
network from Vis. Geom.
Group at Oxford



learning rate, epochs
.00001 phase 1,2
.000001 phase 3,4
5 epochs



Using VGG19 network

detector.Network.Layers(1:7)

- 1 'input' Image Input 224x224x3 images with 'zerocenter' normalization
- 2 'conv1_1' Convolution 64 3x3x3 convolutions with stride [1 1] and padding [1 1 1 1]
- 3 'relu1_1' ReLU ReLU
- 4 'conv1_2' Convolution 64 3x3x64 convolutions with stride [1 1] and padding [1 1 1 1]
- 5 'relu1_2' ReLU ReLU
- 6 'pool1' Max Pooling 2x2 max pooling with stride [2 2] and padding [0 0 0 0]
- 7 'conv2_1' Convolution 128 3x3x64 convolutions with stride [1 1] and padding [1 1 1 1]


.....

detector.Network.Layers(35:47)

- 1 'relu5_3' ReLU ReLU
- 2 'conv5_4' Convolution 512 3x3x512 convolutions with stride [1 1] and padding [1 1 1 1]
- 3 'relu5_4' ReLU ReLU
- 4 'roi pooling layer' ROI Max Pooling ROI Max Pooling with grid size [7 7]
- 5 'fc6' Fully Connected 4096 fully connected layer
- 6 'relu6' ReLU ReLU
- 7 'drop6' Dropout 50% dropout
- 8 'fc7' Fully Connected 4096 fully connected layer
- 9 'relu7' ReLU ReLU
- 10 'drop7' Dropout 50% dropout
- 11 'fc_detection' Fully Connected 2 fully connected layer
- 12 'softmax' Softmax softmax
- 13 'classoutput' Classification Output crossentropyex with classes 'hand' and 'Background'

Results so far

- Using 50% overlap with true box as correct ~50% TP rate
- 4-8 hours on 1 compute node (CPU) 5 epochs on ~1K images



In Summary

faster RCNN get near 60% TP rate

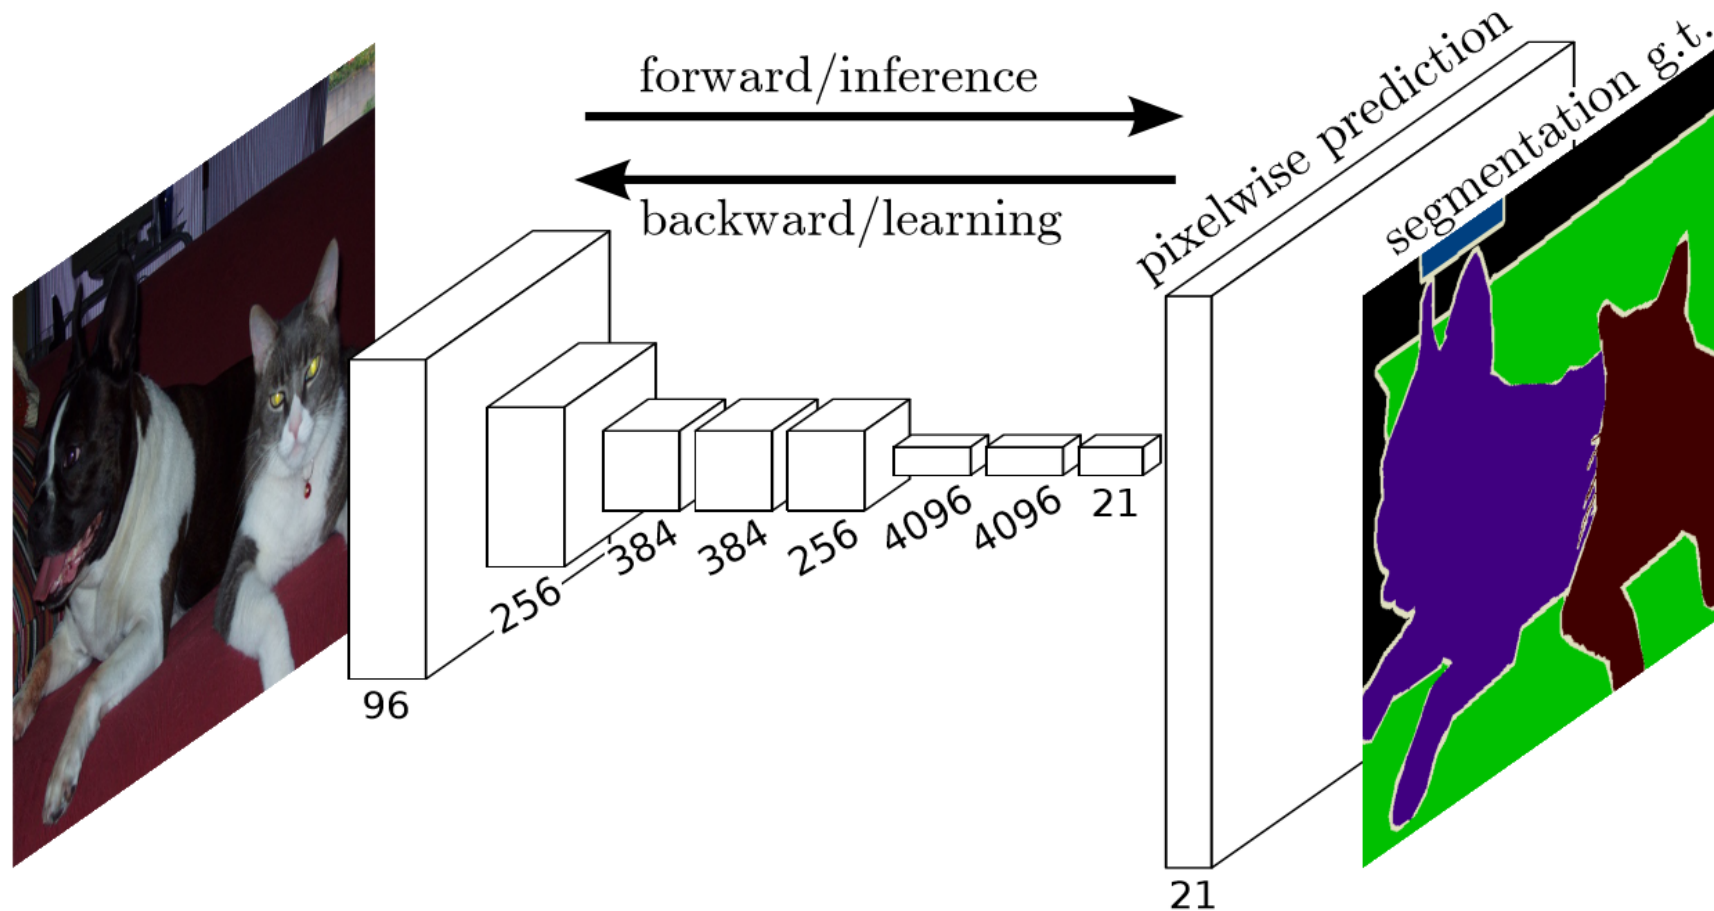
(<https://pjreddie.com/media/files/papers/YOLOv3.pdf>)

proposing regions takes much time

Matlab fasterRCNN easy to use but could use more options


(like turning off weight changes for transfer RCNN learning, reading in images faster, etc..)

Learning Segmentation (deconvolve)



Caffe2, Facebook “Detectron” networks

Object Detection
ie getting a region
bounding box
(rcnn)



Object
Segmentation
ie getting a mask
(mask-rcnn)



Object Parts
ie getting keypoints
(keypoint-rcnn)



Caffe2 quick overview

- Facebook took over Caffe, and built Caffe2 on top of pyTorch
- Keras is easier to learn, Caffe2 better for production (supposedly)
- CNNs are built as defined-nets (ie network configurations)
- CNNs are run as prediction-nets
- Network activity directly available as “blobs” (like tensors)
- Caffe2 ‘brew’ library has Keras-like higher level API

Caffe2 Detectron on Comet

- **git clone *<https://github.com/facebookresearch/Detectron>***
You will get folders of tools, utilities, etc..
- **On Comet compute node, run:**
module purge
module load singularity
singularity shell /share/apps/gpu/singularity/images/pytorch/pytorch-v1.0.0-gpu-20190110.simg

Detectron sample execution

“infer_simple” is python program to load and run network and output visualizations

```
python tools/infer_simple.py \
  --cfg configs/12_2017_baselines/e2e_keypoint_rcnn_R-101-
  FPN_1x.yaml \
  --output-dir ./my_output_results \
  --image-ext png \
  --kp-thresh 2 \
  --wts
https://dl.fbaipublicfiles.com/detectron/37697946/12\_2017\_baselines/
e2e\_keypoint\_rcnn\_R-101-
FPN\_1x.yaml.08\_45\_06.Y14KqbST/output/train/keypoints\_coco\_201
4\_train:keypoints\_coco\_2014\_valminusminival/generalized\_rcnn/mod
el\_final.pkl \
  my_directory_of_images_to_input
```

A network configuration file; see github site for list - they vary by size, training sets, etc.. and what they output.

Output keypoints above a score (look in utils/vis.py to see how all output details are returned)

Weights for this network

References

- **Book:** <https://mitpress.mit.edu/books/deep-learning>
- **Documentation:** <https://keras.io/>
- **Tutorials I used (borrowed):**
 - <http://cs231n.github.io/convolutional-networks/>
 - <https://hackernoon.com/visualizing-parts-of-convolutional-neural-networks-using-keras-and-cats-5cc01b214e59>
 - https://github.com/julienr/ipynb_playground/blob/master/keras/convmnist/keras_cnn_mnist.ipynb